

ASSIGNMENT 3

Vaishnav Mule
A20516627

1 Recitation Exercises

1.1 Chapter 6

Exercise: 1

a) In backward selection we can select any term to delete at any phase, But in forward selection we can observe only the immediate pair of terms at a given step. Hence, Backward selection is used to produce lowest training RSS.

b) It is hard to surely pick an option, the model that has the best subset approach will choose the model with k predictors that is best.

c)

1. True
 2. True
 3. False
 4. False
 5. True
-

Exercise: 2

a) the correct option is (iii), Lasso does a better job than least squares is because of the bias-variance trade-off. Also lasso performs variable selection, which is easier to understand than other techniques like ridge regression. Just at a moderate cost increase in bias lasso approach can reduce variance when the variance of the least squares estimates is sufficiently big.

b) option (iii) is correct. Ridge regression is less adaptable than least squares, as the variance falls, and the bias rises as the coefficient nears 0.

c) option (i) is correct, as non-linear regression models are comparatively more flexible.

Exercise: 3

a) option (i) is correct as beta variables have less constraint as s increases as they can now accept values form wider range. Model now fits the training data better.

- b) option (ii) is correct as though the model can initially fit data well, after a point it tends to overfit and the model becomes very flexible.
 - c) option (iii) is correct as more betas are included, the model becomes flexible and there will be a steady increase in variance.
 - d) option (iv) is correct, it is self-explanatory.
 - e) option (v) is correct as error term has nothing to do with the model.
-

Exercise: 4

- a) option (iii) is correct, an increase in λ will minimize beta squared term. Hence, overall fit will get worse for training as λ is increased.
 - b) option (i) is correct, the test error will initially decrease as variance drops but the test MSE will grow if λ increases as a result of variance increasing.
 - c) option (iv) is correct, several beta values are very close to zero hence the complexity of the model will decrease on the whole.
 - d) option (iii) is correct as model gets simpler.
 - e) option (v) is correct as it can't be controlled. Error term is not dependant on the model.
-

Exercise: 5

a)

① 5 a)
$$L = (y_1 - \hat{\beta}_1 x_{11} - \hat{\beta}_2 x_{12})^2 + (y_2 - \hat{\beta}_1 x_{21} - \hat{\beta}_2 x_{22})^2 + \lambda (\hat{\beta}_1^2 + \hat{\beta}_2^2)$$

 let $x_{11} = x_{21} = x_1, x_{12} = x_{22} = x_2$

$$L = (y_1 - (\hat{\beta}_1 + \hat{\beta}_2) x_1)^2 + (y_2 - (\hat{\beta}_1 + \hat{\beta}_2) x_2)^2 + \lambda (\hat{\beta}_1^2 + \hat{\beta}_2^2)$$

b)

b) differentiate ① with respect to $\hat{\beta}_1, d y \hat{\beta}_2$

$$\frac{\partial L}{\partial \hat{\beta}_1} = 2(y_1 - (\hat{\beta}_1 + \hat{\beta}_2) x_1)(-x_1) + 2(y_2 - (\hat{\beta}_1 + \hat{\beta}_2) x_2)(-x_2) + 2\lambda \hat{\beta}_1 = 0$$

$$\Rightarrow -x_1 y_1 + (\hat{\beta}_1 + \hat{\beta}_2) x_1^2 + x_2 y_2 + (\hat{\beta}_1 + \hat{\beta}_2) x_2^2 + \lambda \hat{\beta}_1 = 0$$

$$\Rightarrow (x_1^2 + x_2^2 + \lambda) \hat{\beta}_1 + (x_1^2 + x_2^2) \hat{\beta}_2 = x_1 y_1 + x_2 y_2 \quad \text{--- (2)}$$

$$\frac{\partial L}{\partial \hat{\beta}_2} = 2(y_1 - (\hat{\beta}_1 + \hat{\beta}_2) x_1)(-x_1) + 2(y_2 - (\hat{\beta}_1 + \hat{\beta}_2) x_2)(-x_2) + 2\lambda \hat{\beta}_2 = 0$$

$$= -x_1 y_1 + (\hat{\beta}_1 + \hat{\beta}_2) x_1^2 - x_2 y_2 + (\hat{\beta}_1 + \hat{\beta}_2) x_2^2 + \lambda \hat{\beta}_2 = 0$$

$$\Rightarrow \hat{\beta}_2 (x_1^2 + x_2^2 + \lambda) + \hat{\beta}_1 (x_1^2 + x_2^2) = x_1 y_1 + x_2 y_2 \quad \text{--- (3)}$$

Subtract ② from ③

$$\Rightarrow \hat{\beta}_1 (x_1^2 + x_2^2) = \hat{\beta}_1 (x_1^2 + x_2^2 + \lambda) + \hat{\beta}_2 (x_1^2 + x_2^2 + \lambda) - \hat{\beta}_2 (x_1^2 + x_2^2) = 0$$

$$\Rightarrow \hat{\beta}_1 \lambda + \lambda \hat{\beta}_2 = 0$$

$$\hat{\beta}_1 = \hat{\beta}_2$$

c) d)

c)
$$L = (y_1 - (\hat{\beta}_1 + \hat{\beta}_2)x_1)^2 + (y_2 - (\hat{\beta}_1 + \hat{\beta}_2)x_2)^2 + \lambda(|\hat{\beta}_1| + |\hat{\beta}_2|)$$

d) consider optimization problem,

$$(y_1 - (\hat{\beta}_1 + \hat{\beta}_2)x_1)^2 + (y_2 - (\hat{\beta}_1 + \hat{\beta}_2)x_2)^2$$
 subject to $|\hat{\beta}_1| + |\hat{\beta}_2| \leq 5$

using given constraints in given question, we minimize

$$2(y_1 - (\hat{\beta}_1 + \hat{\beta}_2)x_1)^2 > 0$$

possible solution is $y_1 - (\hat{\beta}_1 + \hat{\beta}_2)x_1 = 0$

$$\Rightarrow \hat{\beta}_1 + \hat{\beta}_2 = y_1/x_1$$

we have a set of solutions:-

$$\Rightarrow (\hat{\beta}_1, \hat{\beta}_2) = \hat{\beta}_1 + \hat{\beta}_2 = 5 \text{ with } \hat{\beta}_1, \hat{\beta}_2 \geq 0$$


$$\Rightarrow (\hat{\beta}_1 + \hat{\beta}_2) = -3 \text{ with } \hat{\beta}_1, \hat{\beta}_2 < 0$$

1.2 Chapter 7


Exercise: 2

- a) Lambda approaching infinity indicates that the first term is not relevant. When $m=0$, $g(0)=g$ and therefore $g' = 0$
- b) Lambda approaching infinity indicates that the first term will not be relevant. When $m = 0$, $g(1) = g'$. Horizontal line would denote function \hat{g} which minimizes first derivative.
- c) Lambda approaching infinity indicates that the first term will not be relevant. When $m = 0$, $g(2) = g''$. A linear line would denote function \hat{g} which minimizes the second derivative.
- d) Lambda approaching infinity indicates that the first term will not be relevant. When $m = 0$, $g(3) = g'''$. The functions \hat{g} that minimizes the second derivative is a quadratic.
- e) Lambda = 0 indicates that the second term won't be relevant. Equation then becomes std least squares eq.

Exercise: 3



Exercise: 4



Exercise: 5

- a) The \hat{g}_2 will have a smaller RSS as it will be more flexible, due to the order of penalty term.
- b) the \hat{g}_1 will have smaller test RSS if \hat{g}_2 overfits because of the higher flexibility.
- c) \hat{g}_1 and \hat{g}_2 will equal each other when $\lambda = 0$.

2 Practicum Problems

2.1 Problem 1

Vaishnavi

2023-10-05

```
install.packages("doMC", repos="http://R-Forge.R-project.org")
## Installing package into 'C:/Users/91951/AppData/Local/R/win-library/4.2'
## (as 'lib' is unspecified)
## installing the source package 'doMC'
```

loading data :

```
mtcars_df = data.frame(mtcars)
head(mtcars_df)

##           mpg  cyl  disp  hp  drat    wt   qsec  vs  am  gear  carb
## Mazda RX4      21.0   6  160  110  3.90  2.620  16.46  0  1    4    4
## Mazda RX4 Wag   21.0   6  160  110  3.90  2.875  17.02  0  1    4    4
## Datsun 710      22.8   4  108   93  3.85  2.320  18.61  1  1    4    1
## Hornet 4 Drive   21.4   6  258  110  3.08  3.215  19.44  1  0    3    1
## Hornet Sportabout 18.7   8  360  175  3.15  3.440  17.02  0  0    3    2
## Valiant         18.1   6  225  105  2.76  3.460  20.22  1  0    3    1
```

split data to train set and test set :

```
library(caret)

## Loading required package: ggplot2

## Loading required package: lattice

mtcars_df_split = createDataPartition(mtcars_df$mpg, p = 0.80, list=FALSE)
mtcars_train_set = mtcars_df[mtcars_df_split,]
mtcars_test_set = mtcars_df[-mtcars_df_split,]
```

fitting model :

```
linear_model_mtcars = lm(mpg ~ cyl + disp + hp + drat + wt + qsec + vs + a
m + gear + carb, data = mtcars_train_set)

summary(linear_model_mtcars)
```

```
##
## Call:
## lm(formula = mpg ~ cyl + disp + hp + drat + wt + qsec + vs +
##      am + gear + carb, data = mtcars_train_set)
##
## Residuals:
##      Min        1Q    Median        3Q        Max
## -3.0498 -1.3743  0.0014  0.8413  3.4257
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -35.424905   25.724916  -1.377   0.1864
## cyl           0.637367    1.050158   0.607   0.5519
## disp          0.030854    0.016979   1.817   0.0869 .
## hp           -0.007901    0.020962  -0.377   0.7109
## drat          2.751038    1.626190   1.692   0.1089
## wt           -5.089883    1.775482  -2.867   0.0107 *
## qsec          2.264852    1.029643   2.200   0.0420 *
## vs           -1.192677    2.226921  -0.536   0.5992
## am            1.244666    1.867636   0.666   0.5141
## gear          3.707569    1.687494   2.197   0.0422 *
## carb         -0.552617    0.875073  -0.632   0.5361
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.196 on 17 degrees of freedom
## Multiple R-squared:  0.9041, Adjusted R-squared:  0.8477
## F-statistic: 16.03 on 10 and 17 DF, p-value: 9.257e-07
```

Based on the t-stat, wt seems relevant. Furthermore other features like qsec, hp, disp also have big t-stat value.

-> wt = -4.74715 -> qsec = 2.14918 -> hp = -0.01274 -> disp = 0.02338

```
coef(linear_model_mtcars)
##      (Intercept)          cyl          disp          hp          drat
## -35.424905463    0.637367233    0.030853587  -0.007901284    2.751037674
##           wt          qsec          vs          am          gear
##  -5.089882786    2.264852144  -1.192676997    1.244665660    3.707568754
##           carb
##  -0.552616915
```

performing ridge regression :

```
library(glmnet)

## Loading required package: Matrix

## Loaded glmnet 4.1-6

library(doMC)

## Loading required package: foreach

## Loading required package: iterators
```




```
## Loading required package: parallel

y = mtcars_train_set$mpg
x = data.matrix(mtcars_train_set[,c('cyl','disp','hp','drat','wt','qsec','vs','am','gear','carb')])
glm = glmnet(x,y, alpha= 0)
summary(glm)

##           Length Class      Mode
## a0          100   -none-   numeric
## beta        1000 dgCMatrix S4
## df           100   -none-   numeric
## dim           2   -none-   numeric
## lambda        100   -none-   numeric
## dev.ratio     100   -none-   numeric
## nulldev        1   -none-   numeric
## npasses        1   -none-   numeric
## jerr           1   -none-   numeric
## offset         1   -none-   logical
## call           4   -none-   call
## nobs           1   -none-   numeric

plot(glm)
```



```
registerDoMC(core=2)
grid_lambda = 10^seq(10,-2,length=100)
mtcars_cross_val_glm = cv.glmnet(x,y, alpha= 0, lambda = grid_lambda, parallel = TRUE, grouped=FALSE)
summary(mtcars_cross_val_glm)

##           Length Class      Mode
## lambda        100   -none-   numeric
```




```
## cvm      100    -none- numeric
## cvsd     100    -none- numeric
## cvup     100    -none- numeric
## cvlo     100    -none- numeric
## nzero    100    -none- numeric
## call      7     -none- call
## name      1     -none- character
## glmnet.fit 12    elnet  list
## lambda.min 1     -none- numeric
## lambda.1se 1     -none- numeric
## index     2     -none- numeric

best_lambda_val = mtcars_cross_val_glm$lambda.min
best_lambda_val

## [1] 2.656088

plot(mtcars_cross_val_glm)
```



```
mtcars_new_model = glmnet(x, y, alpha=0, lambda=best_lambda_val)
summary(mtcars_new_model)
```

##	Length	Class	Mode
## a0	1	-none-	numeric
## beta	10	dgCMatrix	S4
## df	1	-none-	numeric
## dim	2	-none-	numeric
## lambda	1	-none-	numeric
## dev.ratio	1	-none-	numeric
## nulldev	1	-none-	numeric
## npasses	1	-none-	numeric
## jerr	1	-none-	numeric

```
## offset      1      -none-    logical
## call        5      -none-    call
## nobs         1      -none-    numeric

y_test = mtcars_test_set$mpg
x_test = data.matrix(mtcars_test_set[,c('cyl','disp','hp','drat','wt','qsec','vs','am','gear','carb')])
y_mtcars_predictions = predict(mtcars_new_model, s=best_lambda_val, newx=x_test)
```

results Sample Test :

```
SST_mtcars = sum((y_test - mean(y_test))^2)
cat("SST : ", SST_mtcars)

## SST : 270.91

SSE_mtcars = sum((y_mtcars_predictions - mean(y_test))^2)
cat("\nSSE : ", SSE_mtcars)

##
## SSE : 109.7211

mtcars_R_sq = 1 - (SSE_mtcars/SST_mtcars)
cat("\nR -square : ", mtcars_R_sq)

##
## R -square : 0.5949908
```

Difference in coefficient values of models can be seen below :

```
coef(mtcars_new_model)


## 11 x 1 sparse Matrix of class "dgCMatrix"
##              s0
## (Intercept) 18.109297624
## cyl         -0.326445530
## disp        -0.003237479
## hp          -0.010148915
## drat         1.227817848
## wt          -1.242302378
## qsec         0.192948571
## vs           0.636884201
## am           1.441677967
## gear         0.816995641
## carb        -0.525317638

coef(linear_model_mtcars)

##      (Intercept)      cyl      disp      hp      drat
## -35.424905463    0.637367233    0.030853587 -0.007901284    2.751037674
##      wt      qsec      vs      am      gear
## -5.089882786    2.264852144 -1.192676997    1.244665660    3.707568754
##      carb
## -0.552616915
```

It can be seen from new coefficient values that Ridge Regression did not do variable selection (coeff = 0). It performed 'Shrinkage'.

```
plot(glm, xvar="lambda", label=TRUE)
```



2.2 Problem 2

Vaishnavi

2023-10-05

Load Dataset :

```
swiss_df = data.frame(swiss)
head(swiss)
```

```
##           Fertility Agriculture Examination Education Catholic
## Courtelary      80.2         17.0           15         12      9.96
## Delemont        83.1         45.1            6          9     84.84
## Franches-Mnt    92.5         39.7            5          5     93.40
## Moutier         85.8         36.5           12          7     33.77
## Neuveville      76.9         43.5           17         15      5.16
## Porrentruy      76.1         35.3            9          7     90.57
##           Infant.Mortality
## Courtelary             22.2
## Delemont               22.2
## Franches-Mnt           20.2
## Moutier                20.3
## Neuveville             20.6
## Porrentruy             26.6
```

```
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```

swiss_df_split = createDataPartition(swiss_df$Fertility, p = 0.80, list=FALSE)
swiss_df_train_set = swiss_df[swiss_df_split,]
swiss_df_test_set = swiss_df[-swiss_df_split,]

```

Fit a linear model :

```

swiss_linearmodel = lm(Fertility ~ Agriculture + Examination + Education +
Catholic + Infant.Mortality, data = swiss_df_train_set)
summary(swiss_linearmodel)

```

```

##
## Call:
## lm(formula = Fertility ~ Agriculture + Examination + Education +
##      Catholic + Infant.Mortality, data = swiss_df_train_set)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -17.0815  -4.2016  -0.1039   4.5836  13.4483
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   72.54340    12.66230   5.729 2.14e-06 ***
## Agriculture   -0.16029     0.07854  -2.041  0.04934 *
## Examination   -0.42089     0.32496  -1.295  0.20425
## Education     -0.72175     0.23333  -3.093  0.00401 **
## Catholic       0.09250     0.04403   2.101  0.04340 *
## Infant.Mortality 0.85870     0.42425   2.024  0.05112 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 7.416 on 33 degrees of freedom
## Multiple R-squared:  0.6973, Adjusted R-squared:  0.6514
## F-statistic: 15.2 on 5 and 33 DF, p-value: 9.292e-08

```

Based on t-values and p-values, probable relevant features are Infant. Mortality and education

coefficient (education) = -0.83356522 coefficient (Infant.Mortality) = 1.20966304

```

coef(swiss_linearmodel)
##      (Intercept)      Agriculture      Examination      Education
##      72.54340033      -0.16028654      -0.42088540      -0.72174890
##      Catholic Infant.Mortality
##      0.09249552      0.85870229

```

Perform Lasso Regression

```

library(glmnet)

## Loading required package: Matrix

## Loaded glmnet 4.1-6

library(doMC)


```

```
## Loading required package: foreach
## Loading required package: iterators
## Loading required package: parallel

swiss_x_train = data.matrix(swiss_df_train_set[,c("Agriculture", "Examination", "Education", "Infant.Mortality")])
swiss_y_train = swiss_df_train_set$Fertility
swiss_lassomodel = glmnet(swiss_x_train, swiss_y_train, alpha=1)
summary(swiss_lassomodel)

##           Length Class      Mode
## a0           65    -none-   numeric
## beta        260 dgCMatrx S4
## df           65    -none-   numeric
## dim           2    -none-   numeric
## lambda        65    -none-   numeric
## dev.ratio     65    -none-   numeric
## nulldev        1    -none-   numeric
## npasses        1    -none-   numeric
## jerr           1    -none-   numeric
## offset         1    -none-   logical
## call           4    -none-   call
## nobs           1    -none-   numeric

plot(swiss_lassomodel)
```



```
registerDoMC(cores=2)
lambda_vec = 10^seq(10, -3, length=100)
swiss_crossval_model = cv.glmnet(swiss_x_train, swiss_y_train, alpha=1, la
```


```
mbda=lambda_vec, parallel=TRUE, grouped=FALSE)
summary(swiss_crossval_model)

##           Length Class   Mode
## lambda      100    -none-  numeric
## cvm          100    -none-  numeric
## cvsd         100    -none-  numeric
## cvup         100    -none-  numeric
## cvlo         100    -none-  numeric
## nzero        100    -none-  numeric
## call          7    -none-   call
## name          1    -none- character
## glmnet.fit    12    elnet   list
## lambda.min     1    -none-  numeric
## lambda.1se     1    -none-  numeric
## index          2    -none-  numeric

swiss_bestlambda_val = swiss_crossval_model$lambda.min
swiss_bestlambda_val

## [1] 0.001

plot(swiss_crossval_model)
```



```
swiss_new_model = glmnet(swiss_x_train, swiss_y_train, alpha=1, lambda=swi
ss_bestlambda_val)
summary(swiss_new_model)

##           Length Class      Mode
## a0          1    -none-   numeric
## beta         4    dgMatrix S4
## df           1    -none-   numeric
```

```

## dim      2      -none-    numeric
## lambda   1      -none-    numeric
## dev.ratio 1      -none-    numeric
## nulldev  1      -none-    numeric
## npasses  1      -none-    numeric
## jerr      1      -none-    numeric
## offset   1      -none-    logical
## call      5      -none-    call
## nobs      1      -none-    numeric

coef(swiss_new_model)

## 5 x 1 sparse Matrix of class "dgCMatrix"
##              s0
## (Intercept)  79.5714706
## Agriculture -0.1336910
## Examination -0.8928483
## Education   -0.4621702
## Infant.Mortality 0.8808018

coef(swiss_linearmodel)

##      (Intercept)      Agriculture      Examination      Education
##      72.54340033      -0.16028654      -0.42088540      -0.72174890
##      Catholic Infant.Mortality
##      0.09249552      0.85870229

swiss_x_test = data.matrix(swiss_df_test_set[,c("Agriculture", "Examination", "Education", "Infant.Mortality")])
swiss_y_test = swiss_df_test_set$Fertility
swiss_y_predictions = predict(swiss_new_model, newx=swiss_x_test, lambda=swiss_bestlambda_val)

SST_swiss = sum((swiss_y_test - mean(swiss_y_test))^2)
cat("SST : ", SST_swiss)

## SST :  1175.359

SSE_swiss = sum((swiss_y_predictions - mean(swiss_y_test))^2)
cat("\nSSE : ", SSE_swiss)


##
## SSE :  282.8698

swiss_R_sq = 1 - (SSE_swiss/SST_swiss)
cat("\nR -square : ", swiss_R_sq)

##
## R -square :  0.7593332

plot(swiss_lassomodel, xvar="lambda", label=TRUE)

```

Model selection not performed by Lasso as none of the coefficients are = 0

2.2 Problem 3

Vaishnavi

2023-10-05

Load the Dataset :

```
library(readxl)
concrete_df = read_excel('Concrete_Data.xls')
names(concrete_df) = c('cement', 'blastfurnaceslag', 'fly_ash', 'water', 'superplasticizer', 'coarse_aggregate', 'fine_aggregate', 'age', 'concretecompressstrength')
head(concrete_df)
```

```
## # A tibble: 6 × 9
##   cement blastfurnaceslag fly_ash water superplasticizer coarse_aggregate fine_aggregate age concretecompressstrength
##   <dbl>          <dbl>   <dbl> <dbl>          <dbl>      <dbl>      <dbl> <dbl>
## 1  540            0         0  162            2.5    1040        676    2
## 2  540            0         0  162            2.5    1055        676    2
## 3  332.          142.         0  228            0       932        594   27
```

```

0    40.3
## 4    332.          142.          0    228          0    932    594    36
5    41.1
## 5    199.          132.          0    192          0    978.    826.    36
0    44.3
## 6    266          114          0    228          0    932    670    9
0    47.0
## # ... with abbreviated variable names 1superplasticizer, 2coarse_aggregate,
## # 3fine_aggregate, 4concretecompressstrength

```

Create GAM model :

```

library(mgcv)

## Loading required package: nlme

## This is mgcv 1.8-41. For overview type 'help("mgcv-package")'.

library(nlme)
concrete_gam_model = gam(concretecompressstrength ~ cement + blastfurnaceslag + fly_ash + water + superplasticizer + coarse_aggregate + fine_aggregate + age, data= concrete_df)
summary(concrete_gam_model)

##
## Family: gaussian
## Link function: identity
##
## Formula:
## concretecompressstrength ~ cement + blastfurnaceslag + fly_ash +
##   water + superplasticizer + coarse_aggregate + fine_aggregate +
##   age
##
## Parametric coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -23.163756  26.588421  -0.871 0.383851
## cement        0.119785   0.008489  14.110 < 2e-16 ***
## blastfurnaceslag 0.103847   0.010136  10.245 < 2e-16 ***
## fly_ash       0.087943   0.012585   6.988 5.03e-12 ***
## water        -0.150298   0.040179  -3.741 0.000194 ***
## superplasticizer 0.290687   0.093460   3.110 0.001921 **
## coarse_aggregate 0.018030   0.009394   1.919 0.055227 .
## fine_aggregate 0.020154   0.010703   1.883 0.059968 .
## age          0.114226   0.005427  21.046 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##
## R-sq.(adj) = 0.612   Deviance explained = 61.5%
## GCV = 109.11   Scale est. = 108.16      n = 1030

concrete_gam_model = gam(concretecompressstrength ~ s(cement) + s(blastfurnaceslag) + s(fly_ash) + s(water) + s(superplasticizer) + s(coarse_aggregate) + s(fine_aggregate) + s(age))

```

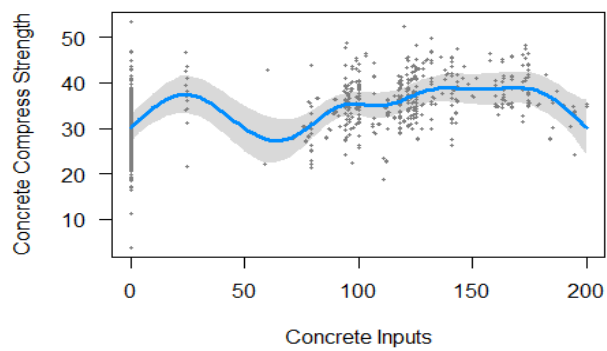

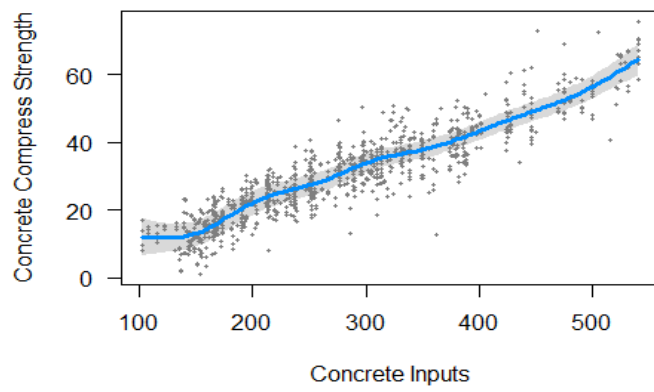
```

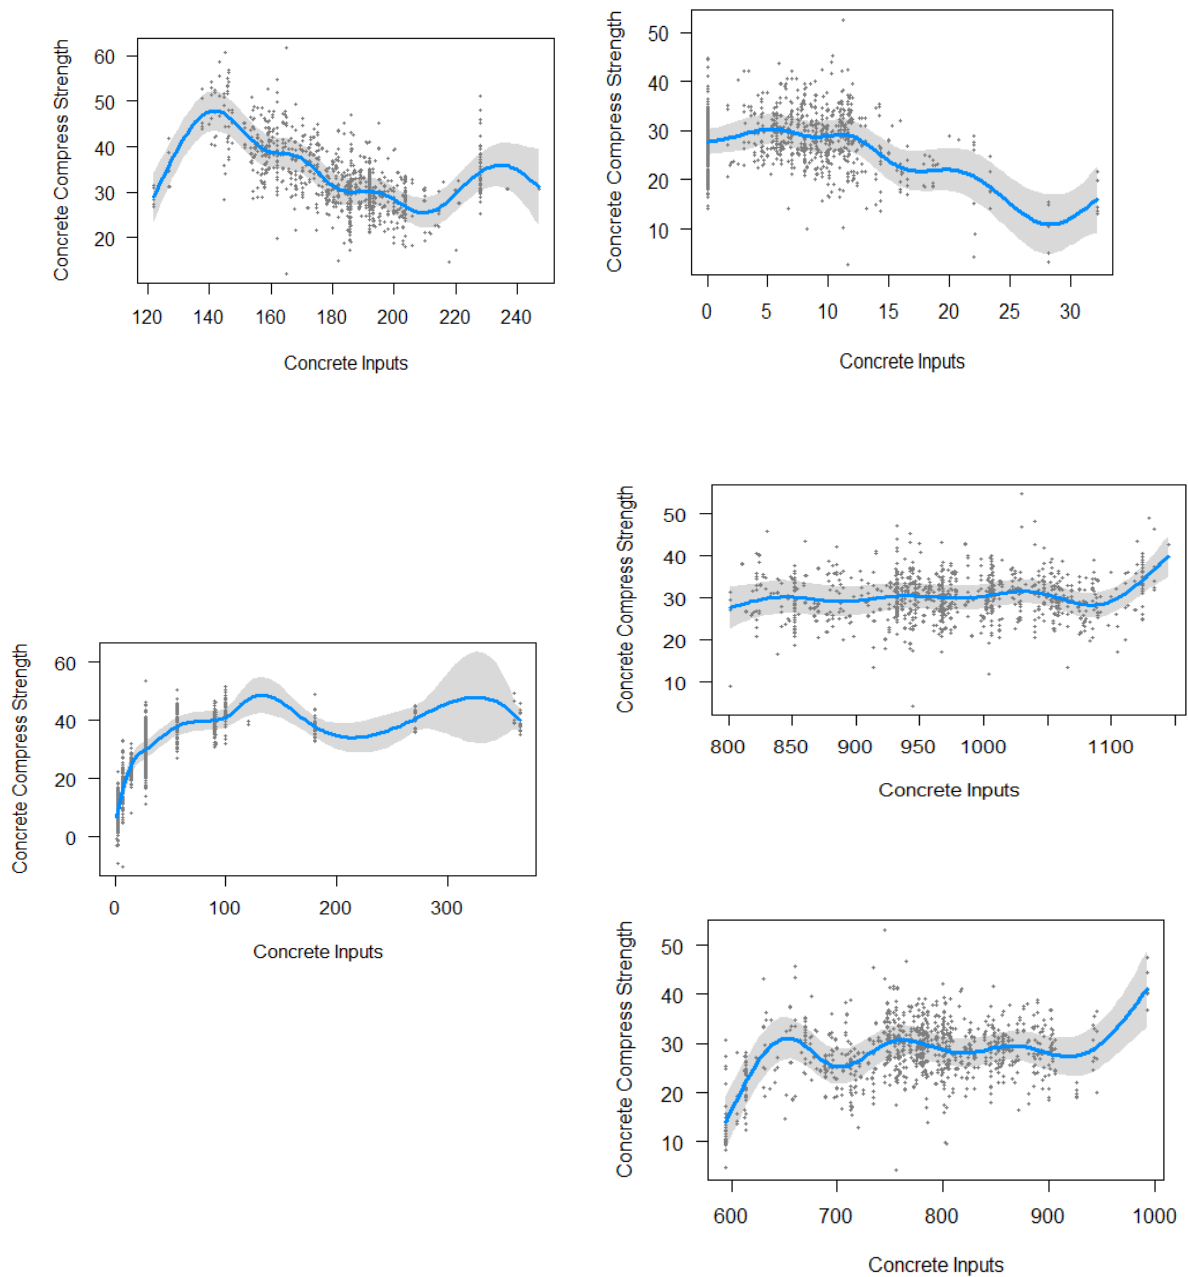
te) + s(fine_aggregate) + s(age), data= concrete_df)
summary(concrete_gam_model)

##
## Family: gaussian
## Link function: identity
##
## Formula:
## concretestrength ~ s(cement) + s(blastfurnaceslag) +
##   s(fly_ash) + s(water) + s(superplasticizer) + s(coarse_aggregate) +
##   s(fine_aggregate) + s(age)
##
## Parametric coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)  35.8178    0.1675   213.9   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##               edf Ref.df      F  p-value
## s(cement)      8.223  8.830 48.690 < 2e-16 ***
## s(blastfurnaceslag) 8.114  8.757 25.041 < 2e-16 ***
## s(fly_ash)      8.256  8.817  9.354 < 2e-16 ***
## s(water)        8.742  8.973 26.116 < 2e-16 ***
## s(superplasticizer) 8.039  8.743 10.845 < 2e-16 ***
## s(coarse_aggregate) 7.904  8.673  3.403 0.000593 ***
## s(fine_aggregate)  8.618  8.951 18.435 < 2e-16 ***
## s(age)          8.559  8.900 365.119 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.896   Deviance explained = 90.3%
## GCV = 30.914   Scale est. = 28.89        n = 1030

library(visreg)
visreg(concrete_gam_model, xlab="Concrete Inputs", ylab = "Concrete Compressive Strength")

```





Although there are certain data points that are not very close to our confidence interval, it is seen that there are predictors that are linearly associated and some non-linearly linked is also quite well fitted.